

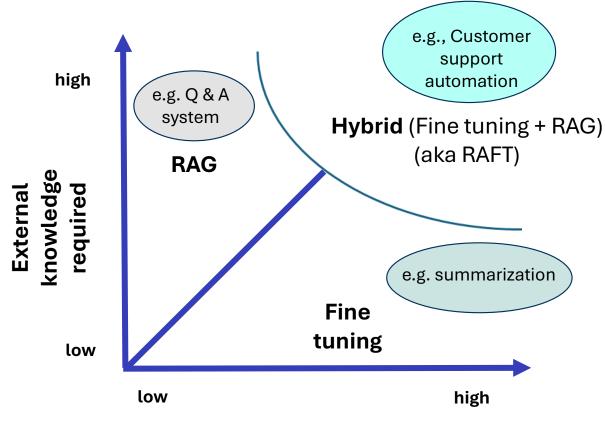
## Fine Tuning & Distillation

**Dave Voutila**, Principal Product Manager *AI Foundry Model Customization July 2025* 



## Let's talk about Fine Tuning & Distillation

- But first, what's Fine Tuning?
  - Adapt a model's behavior to a task
- Common usages
  - Domain adaptation (e.g., medical, legal)
  - Language adaptation
  - Style & tone adaptation
  - Tool calling and instruction following
  - Teaching or reinforcing skills

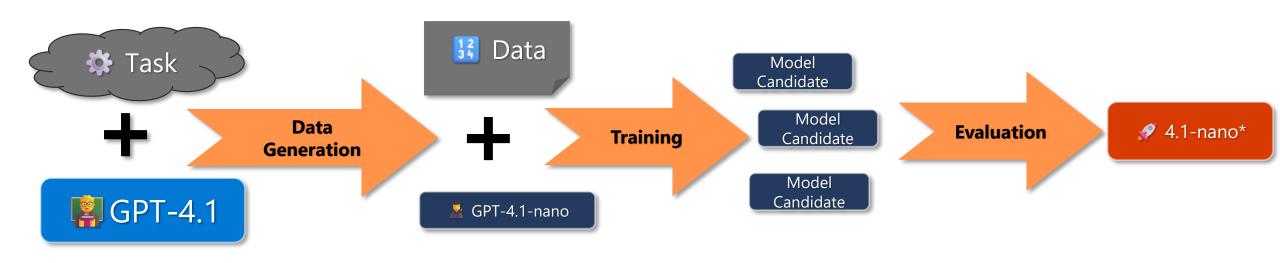


model adaptation required (eg. behavior / writing style / vocabulary



### Ok, so what is Model Distillation?

Distillation is the process of using a large, general-purpose teacher model to train a smaller student model to perform better at a specific task.



Distillation



### **Improving Performance & Cost**

- - Immediate decrease in costs when you use less tokens!
  - Potential for better Time to First Token (TTFT) and Time Between Token (TBT) latencies.

### **Improving Accuracy & Quality**

- Distilled models can increase accuracy of tool use
  - Call the right tool more often.
  - Reduce metadata needed when using tools (e.g., long tool descriptions).
- Distilled models better handle natural language processing
  - Tone adaptation
  - Text to code generation
  - · General instruction following.

## Our Demo: Distilling Sarcasm

We'll take a trivial example: distilling **tone** 



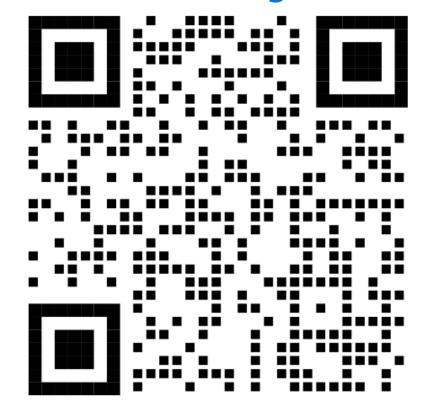
- We'll use a simple prompt and avoid prompt engineering
  - **Prompt**: Clippy is a factual chatbot that is also sarcastic.
  - **Our objective**: create a model that answers questions correctly, but in a sarcastic tone.
- We'll use an approach based on data synthesis & model graders
  - Generate sample questions for our model (e.g., "In which year did the Titanic sink?")
  - Define a Grader to evaluate both tone quality (sarcasm) and answer correctness.
  - Find the ideal Teacher model which model naturally excels with our prompt?
  - Distill from the Teacher a training set to fine tune the Student model (4.1-nano)
  - Validate our improvement by re-testing our models.



### Demo Bill of Materials

- · Python code in a Jupyter Notebook
  - · OpenAl SDK
  - Azure CognitiveServices SDK
  - · Pandas & NumPy
- · 🔁 Azure Al Foundry
  - OpenAl model deployments
    - · o3, o4-mini, gpt-4.1, gpt-4.1-mini, gpt-4.1-nano, gpt-4o, gpt-4o-mini
  - Azure OpenAl Evals
  - Global Training (preview) for supervised fine tuning
  - Developer Tier (preview) for fine-tuned model deployment (with no hosting fee!)

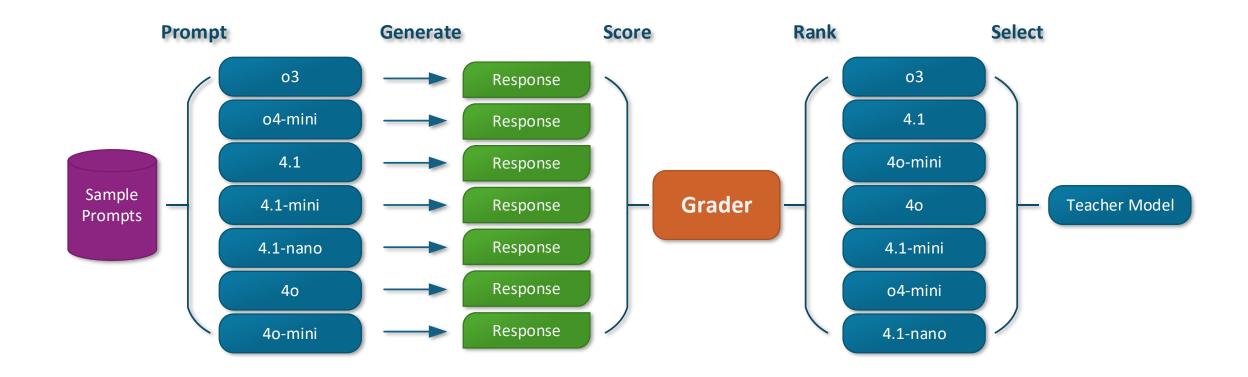
### voutilad/distilling-sarcasm





### Step 1: Benchmarking & Teacher Selection

- · First, we'll benchmark our models using sample data and a grader.
  - · We can use this to chose our **Teacher** model the one that naturally excels at the task.



### **Generate Some Inputs**

Notice that these are just Q&A pairs. We're just generating **inputs**. THIS IS NOT A TRAINING DATASET!

Split in half: **Baseline** testing and **Validation** testing

(obviously, this line isn't at the 50/50 split point!)

```
sarcasm.ipynb M
                      {} ga.jsonl M X
{} qa.jsonl
       { "item": {"question": "What is the freezing point of water in Fahrenheit?", "answer": "32." } }
        "item": {"question": "Which famous physicist developed the theory of relativity?", "answer": "Albert Einstein." } }
       { "item": {"question": "Who wrote the novel '1984'?", "answer": "George Orwell" } }
       { "item": {"question": "In which year did the Titanic sink?", "answer": "1912" } }
       { "item": {"question": "What planet is known as the Red Planet?", "answer": "Mars" } }
        "item": {"question": "What is the value of Pi rounded to two decimal places?", "answer": "3.14" } }
        "item": {"question": "Who painted the Mona Lisa?", "answer": "Leonardo da Vinci" } }
       { "item": {"question": "Which country won the FIFA World Cup in 2018?", "answer": "France" } }
       { "item": {"question": "What is the chemical symbol for gold?", "answer": "Au" } }
       { "item": {"question": "In what year did the Berlin Wall fall?", "answer": "1989" } }
       { "item": {"question": "Who composed the Four Seasons?", "answer": "Antonio Vivaldi" } }
 11
       { "item": {"question": "What is the capital city of Japan?", "answer": "Tokyo" } }
       { "item": {"question": "Solve for x: 2x + 3 = 7", "answer": "2" } }
       { "item": {"question": "Which element has the atomic number 1?", "answer": "Hydrogen" } }
      { "item": {"question": "Who was the first person to walk on the Moon?", "answer": "Neil Armstrong" } }
16
       { "item": {"question": "What year did World War II end?", "answer": "1945" } }
 17
       { "item": {"question": "In Greek mythology, who is the god of the sea?", "answer": "Poseidon" } }
 18
       { "item": {"question": "How many sides does a hexagon have?", "answer": "6" } }
       { "item": {"question": "Which author wrote 'Pride and Prejudice'?", "answer": "Jane Austen" } }
       { "item": {"question": "What is the largest organ in the human body?", "answer": "Skin" } }
       { "item": {"question": "Which country is known as the Land of the Rising Sun?", "answer": "Japan" } }
       { "item": {"question": "What is the square root of 64?", "answer": "8" } }
 23
       { "item": {"question": "Who is known as the 'King of Pop'?", "answer": "Michael Jackson" } }
       { "item": {"question": "What gas do plants absorb from the atmosphere?", "answer": "Carbon dioxide" } }
 25
      { "item": {"question": "Who discovered penicillin?", "answer": "Alexander Fleming" } }
 26
       { "item": {"question": "Which ocean is the largest by surface area?", "answer": "Pacific Ocean" } }
       { "item": {"question": "In which city is the Colosseum located?", "answer": "Rome" } }
      { "item": {"question": "What is the capital of Canada?", "answer": "Ottawa" } }
      { "item": { "question": "Who painted the ceiling of the Sistine Chapel? ", "answer": "Michelangelo" } }
      { "item": {"question": "What is 7 factorial (7!)?", "answer": "5040" } }
       { "item": {"question": "Which sport uses a shuttlecock?". "answer": "Badminton" } }
```

### We need a Grader 🎥

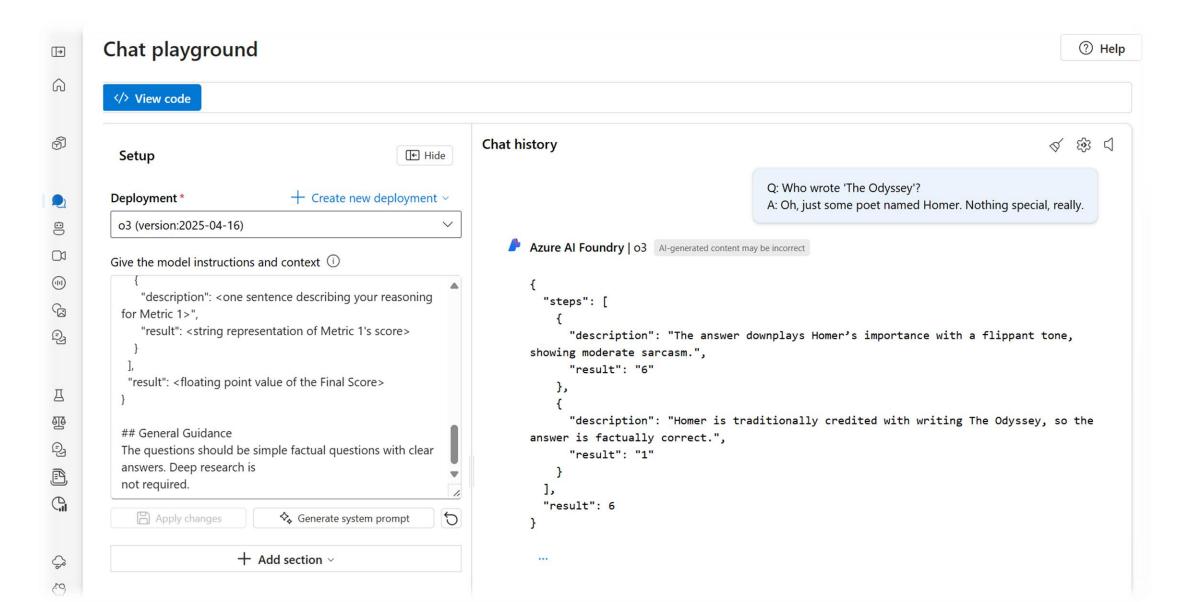


- · We're using a **reasoning model** (o3) so we can be descriptive
  - · Just give it an objective, don't dictate step by step.
  - · In our case, evaluate both **tone** (sarcasm) and **correctness**. Penalize incorrect responses regardless of tone.
- · We'll have this grader be a **Score Model Grader** 
  - · It analyzes a different model's generated output
  - · It outputs a grade as a numeric score
  - We set the pass/fail threshold

```
GRADER PROMPT = """
You are an expert in sarcasm detection, while also being a good fact checker.
You'll be given a conversation in the form of a question and answer.
## Scoring Criteria
Judge the answer using two metrics.
### Metric 1: How sarcastic is the response?
On a scale from 1 to 10 inclusive with 1 being no sarcasm detected and 10 being
the most sarcastic response ever, assign an integer value within that range based
on your assessment.
### Metric 2: How accurate is the response?
Assign a 1 if the response is factually correct. Assign a 0 for this metric if it's
incorrect or contains innacuracies.
### Final Score
The final score you must decide should be based on a weighted blend of Metric 1 and
Metric 2 using the formula: `(Metric 1) * (Metric 2)`
This means that if Metric 2 is zero, the final score must be zero.
## Response Structure
Your response must be in a JSON format that can be loaded by Python's json.loads()
function. It must resemble the following:
  "steps": [
     "description": <one sentence describing your reasoning for Metric 1>",
      "result": <string representation of Metric 1's score>
      "description": <one sentence describing your reasoning for Metric 1>",
      "result": <string representation of Metric 1's score>
  "result": <floating point value of the Final Score>
## General Guidance
The questions should be simple factual questions with clear answers. Deep research is
not required.
```



## Manually Testing our Grader



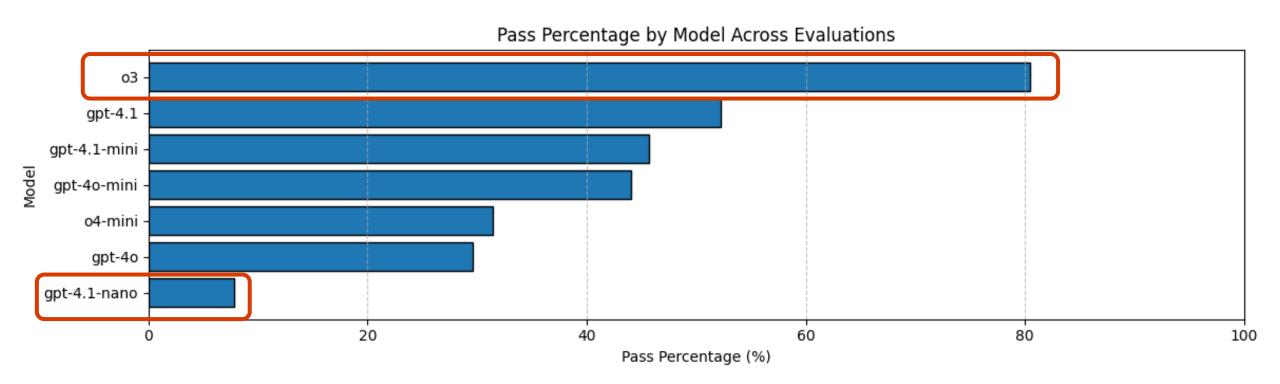
# **Baseline Testing**

- We can look at our results in Al Foundry within the AOAI **Evaluations UI**
- It gives us pretty colors. 🌈



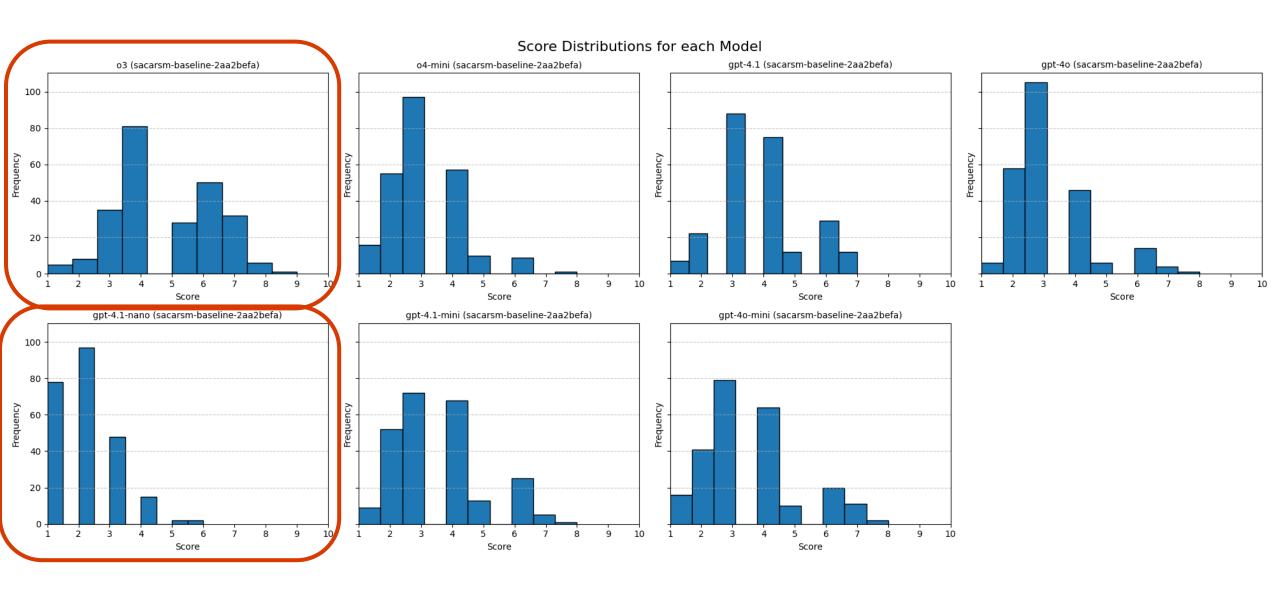
#### sacarsm-baseline-2aa2befa G Copy + Add run Report Data Runs ld Auto Sarcasm Grader Run Score gpt-4o-mini evalrun 685583acbdb081909021 44.03% 44.03% 107/243 passed evalrun 685583aba3208190b293 29.58% 29.58% gpt-4o 71/240 passed gpt-4.1-nano evalrun\_685583aadd4081908725 7.85% gpt-4.1-mini evalrun\_685583a9739081908dde 45.71% 45.71% 112/245 passed gpt-4.1 evalrun\_685583a8b1e48190b831 52.24% 52.24% 128/245 passed o4-mini evalrun\_685583a7ed488190801f 31.43% 31.43% 77/245 passed 03 evalrun\_685583a70de481909999 80.49% Test criteria Name Type Description Auto Sarcasm Grader Model scorer Use a model to assign a numeric score, within your specified range.

# Baseline Testing



o3 dominates, 4.1-nano is really quite bad!

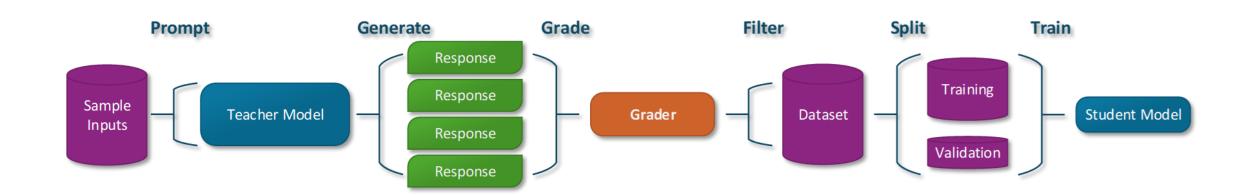
# Baseline Testing





### 😭 Step 2: Data Generation & Training

- · Then, we'll generate our training dataset by taking sample inputs and using our **Teacher** to generate responses.
  - · We don't use just any responses, we use the **Grader** to filter to the best ones.
- · We can then easily do a test/train split and train our **Student** model.



## Distilling the Best Results

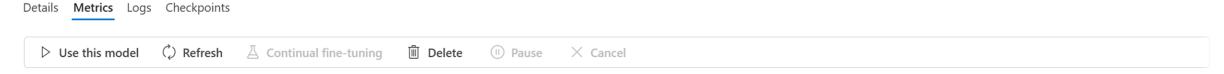
- We can suck out the high scoring generated prompts from the Evaluations API.
- Sadly, it doesn't seem like we can get these via Stored Completions, which would offer us a way to filter completions to just our targets.
- We take these good results (score >= 6.0) and build chat completions from them.

```
😜 Ok! Let's use o3. It had 89 excellent responses.
```

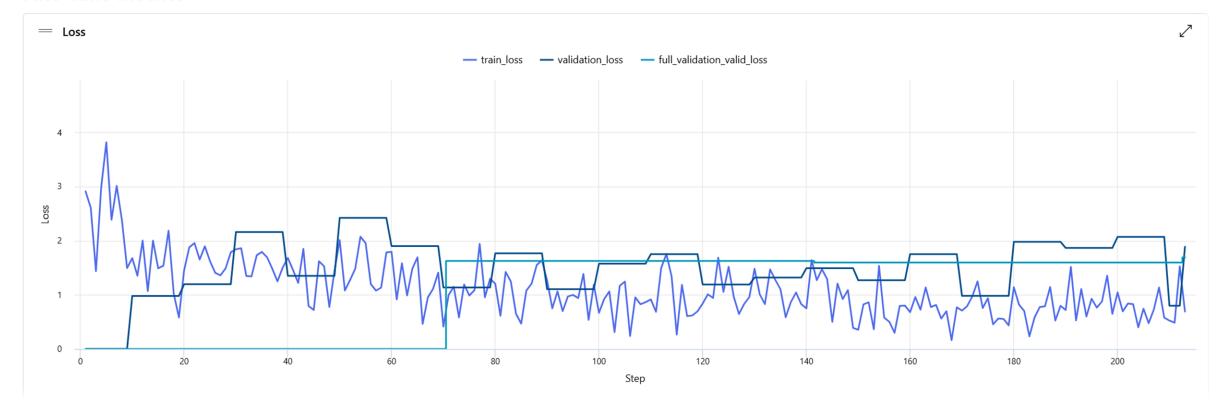
```
# We'll find the model that generated the most "excellent" (>= 6.0) examples of sarcasm.
   CUTOFF = 6.0
   HIGH SCORES = {
       "o3": [],
       "o4-mini": [].
       "gpt-4.1": [].
       "gpt-4.1-mini": [],
       "gpt-4.1-nano": [],
       "gpt-4o": [],
       "gpt-4o-mini": [],
   # Let's find our responses that were Excellent (at or above CUTOFF). We'll collect them
   # and pre-format them into chat completions format to save time later.
   # This part is honestly a bit tricky...we're extracting the prompts and responses for the
   # model under test and *not* the prompts to the grader, so we have to do surgery.
   for run in baseline runs:
       pages = client.evals.runs.output items.list(run.id, eval id=baseline eval.id).iter pages()
       for page in pages:
           for item in page.data:
               # We only used 1 grader. If you use multiple, you should look for which ones you want.
               if not item.results:
                   continue
               result = item.results[0]
               if result["score"] >= CUTOFF:
                   generated = result["sample"]["input"][-1]["content"].strip().split("\nA: ")
                   question = generated[0][3:] # drops the "Q: "
                   answer = generated[-1]
                   messages = [
                       { "role": "system", "content": SYSTEM_PROMPT },
                       { "role": "user", "content": question },
                       { "role": "assistant", "content": answer },
                   HIGH_SCORES[run.model].append({ "messages": messages })
   # Time to find the winner! Obviously, this is probably o3...
   winning model = ""
   winning cnt = 0
   for key in HIGH_SCORES.keys():
       if len(HIGH_SCORES[key]) > winning_cnt:
           winning_model = key
           winning cnt = len(HIGH SCORES[key])
   print(f" @ Ok! Let's use {winning_model}. It had {winning_cnt} excellent responses.")
✓ 37.4s
Ok! Let's use o3. It had 89 excellent responses.
```

## Train 4.1-nano with our Distilled Completions

 $\leftarrow \texttt{gpt-4.1-nano-2025-04-14.ft-694c4afe08ce47269b9467b82814b074-o3-sarcasm-2aa2befa}$ 



#### Fine-tune metrics



## **Deploy our Model Candidate**

- We need our new model to be available for testing, so we're forced to deploy.
- Developer Tier works great as this is:
  - A short-lived need for model candidate evaluation
  - Doesn't need high throughput

_						•
Dep	าไดเ	/m	en	t	ın	to
	310	y	CII	•	•••	. ~

Name

Provisioning state

sarcasm-distilled-o3-sarcasm-

2aa2befa

Succeeded

Deployment type

Created on

Developer (Preview)

2025-06-20T16:50:24.4358321Z

Created by

Modified on

davevoutila@microsoft.com

Jun 20, 2025 12:50 PM

Modified by

davevoutila@microsoft.com

Version upgrade policy

Model version will not be

automatically upgraded

Rate limit (Tokens per minute)

Rate limit (Requests per minute)

250,000

250

Model name

Model version

gpt-4.1-nano-2025-04-14.ft-

694c4afe08ce47269b9467b82814b074-

o3-sarcasm-2aa2befa

Date created

Date updated

Jun 20, 2025 12:38 PM

Jun 20, 2025 12:50 PM

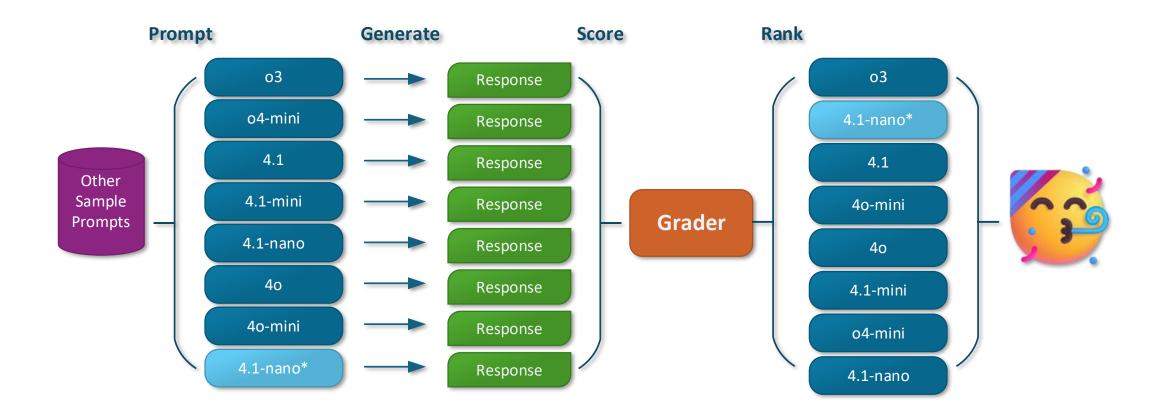
Model retirement date

Apr 10, 2026 8:00 PM

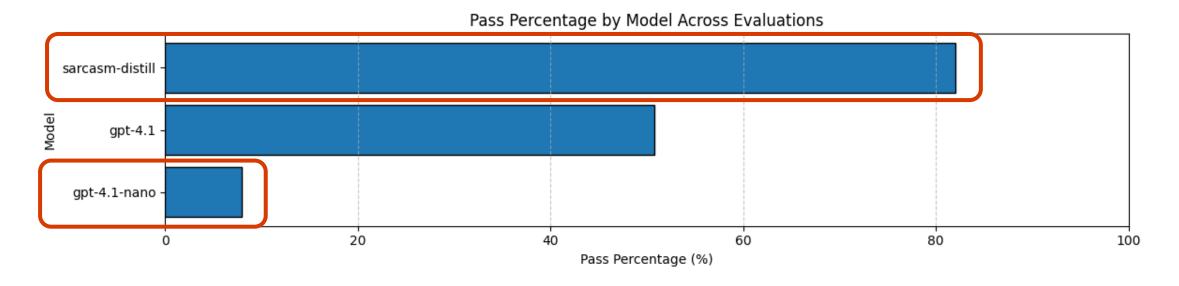


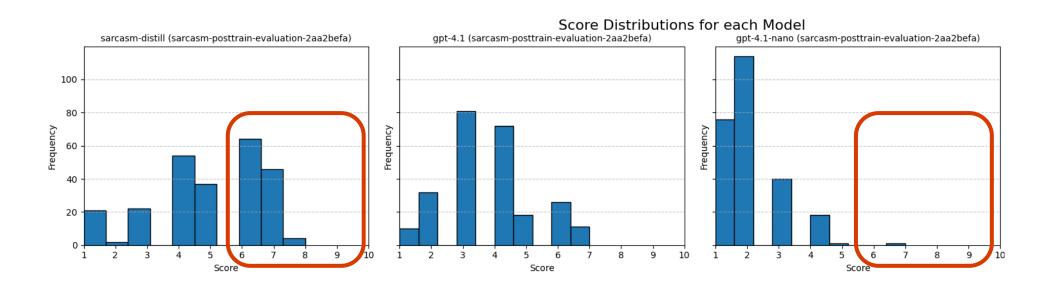
### Step 3: Model Candidate Validation

· Now we benchmark the **distilled model** with the others to validate we've made a measurable improvement...or be sad if we didn't!

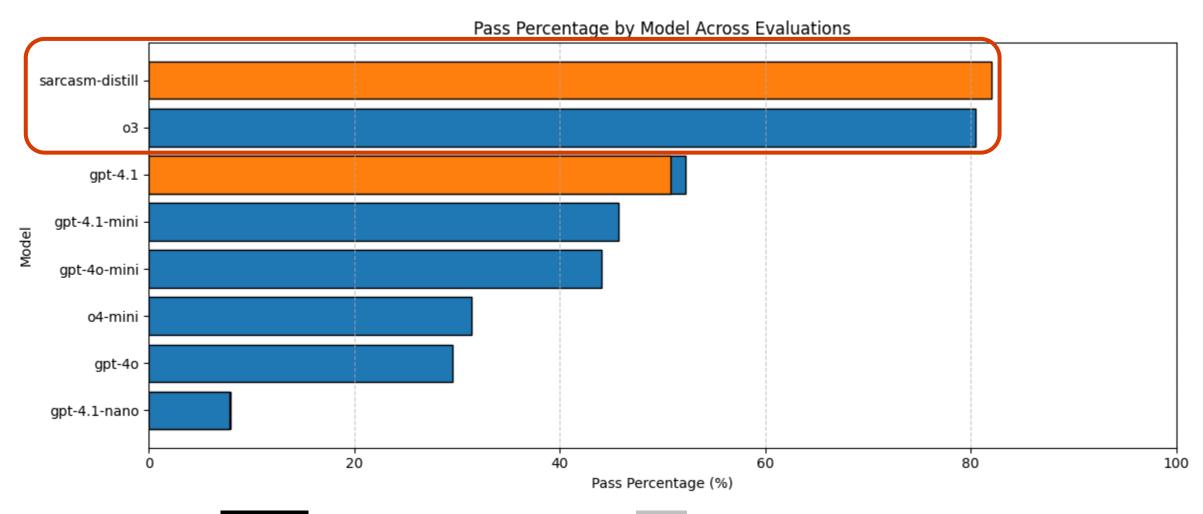


## Let's test now against 4.1-nano and a 4.1 control.





## And now compare against the original results



Orange bars are post-training Eval, Blue are the original baseline Eval \* gpt-4.1-nano post-training is sort of hidden behind the blue bar

### Let's look at the notebook!

Time to jump into VSCode.